

Presenter: Bhanu Teja

Industrial Applications of Advanced Reinforcement Learning





Key Industrial Applications (I)



Smart Manufacturing

Intelligent Production Scheduling

Reinforcement learning optimizes machine and line schedules, cutting idle time through real-time decisions, coordinating tools, workers, robots, and materials for smoother, more efficient production.

Adaptive Control & Maintenance

RL-driven process control tunes parameters like temperature, speed, and pressure, while predictive maintenance selects optimal service timing, increasing throughput, reducing defects, and minimizing unplanned downtime.

Customer Satisfaction	Recommendation rate
8.1	70%
8.2	71%
8.6	76%
7.9	69%
8.0	70%

Autonomous Vehicles



RL-Based Driving Decisions

Reinforcement learning optimizes path planning, lane-changing, and driving behavior including acceleration, braking, and steering to balance safety, comfort, and efficiency in complex traffic environments.



Simulation-to-Reality & Fleet Coordination

Policies are trained in simulation then fine-tuned in controlled real settings; coordinated fleets enable platooning, intersection management, improved energy efficiency, reduced travel time, and enhanced overall road safety.



Healthcare Decision Support in RL



Core healthcare RL application areas

Treatment planning optimizes multi-step dosage and timing, personalized medicine tailors decisions to individual data, ICU support manages ventilation, fluids, drugs under evolving patient conditions.

Constraints and practical implementation

Strict safety, interpretability, and regulatory constraints shape deployment; healthcare reinforcement learning frequently uses offline RL trained on extensive historical patient records.





02

Comparative RL Methods & Summary



RL Method Types

Value vs Policy

Value-based Focus

Value-based methods estimate action values on discrete choices, achieving reliable decisions but often requiring many samples for useful learning.

Discrete Action Setting

These methods suit tasks with fixed action sets, enabling clear comparisons among options but limiting use in rich continuous control domains.

Sample Inefficiency Issue

Learning accurate value estimates usually demands extensive interaction data, making training slow and costly in industrial environments.

VS

Learning vs Models

Model-based Focus

Model-based methods learn environment dynamics to plan, improving sample efficiency but becoming highly sensitive to model inaccuracies.

Efficient Data Usage

By reusing learned models for simulated rollouts, they reduce real-world data needs, valuable when interactions are costly or risky.

Model Error Sensitivity

Bias in the learned dynamics can mislead planning, so validation, uncertainty handling, and robust design are critical in practice.

Thank You



